IS 733 — Lab 11 | Part II: NLP Classification of Data Stories

Paper: *Show and Tell: Exploring LLM's Potential in Formative Educational Assessment of Data Stories*

Dataset: `data_stories_one_shot.csv` — 130 LLM-generated (GPT-4o, one-shot) narrative sentences across 12 data plots

Label: `Show = 0` (Stage 1) | `Tell = 1` (Stage 2 or 3)

In [1]:
```
import re
import warnings
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import roc_auc_score, accuracy_score
from sklearn.preprocessing import MinMaxScaler
from scipy.sparse import hstack, csr_matrix

warnings.filterwarnings('ignore')
np.random.seed(42)
print('Libraries loaded.')
```

```
Libraries loaded.
```

Load & Inspect the Real Dataset

In [2]:
```
# Load dataset
df = pd.read_csv(r'C:\Users\ahmed\Downloads\data_stories_one_shot.csv')

print(f'Shape   : {df.shape}')
print(f'Columns : {df.columns.tolist()}')
print()
print(df.head(8).to_string(index=False))
```

```
Shape   : (130, 4)
Columns : ['Plot_Name', 'Stage', 'Quality', 'Sentence']

Plot_Name  Stage  Quality
Sentence
 walk dog      1      1.0
This is a line chart with error bars.
 walk dog      1      1.0
The chart title is 'Walk dog'.
 walk dog      1      1.0
The y-axis represents 'Mean anxiety'.
 walk dog      1      1.0                The x-axis indicates conditions such as 'Ba
seline', 'After 30 minutes', and presence or absence of 'Dog'.
 walk dog      1      1.0                                                   The chart
compares mean anxiety levels with and without a dog over time.
 walk dog      2      1.0                The mean anxiety level without a dog starts
at around 6 and remains close to this value after 30 minutes.
 walk dog      2      1.0                The mean anxiety level with a dog starts at
around 6 and drops significantly to below 4 after 30 minutes.
 walk dog      2      1.0 There is a downward trend in anxiety levels for both condi
tions, but the drop is more pronounced in the 'Dog' condition.
```

In [3]:

```python
print('=== Dataset Summary ===')
print(f'Total sentences : {len(df)}')
print(f'Unique plots    : {df["Plot_Name"].nunique()}')
print()
print('Stage distribution (raw):')
print(df['Stage'].value_counts().sort_index())
print()
print('Quality distribution:')
print(df['Quality'].value_counts().sort_index())
print()
print('Sentences per plot:')
print(df['Plot_Name'].value_counts().to_string())
```

```
=== Dataset Summary ===
Total sentences : 130
Unique plots    : 12

Stage distribution (raw):
Stage
1    79
2    48
3     3
Name: count, dtype: int64

Quality distribution:
Quality
0.0      9
0.5      6
1.0    115
Name: count, dtype: int64

Sentences per plot:
Plot_Name
tropic storm      13
wealth gap        12
youtube           12
vaccine rate      12
walk dog          11
meaningful        11
vaccine           11
MAP               10
time use          10
solar activity    10
STEM               9
degree             9
```

Create Binary Label

Per the paper: **Show = Stage 1 → label 0 | Tell = Stage 2 or 3 → label 1**

In [4]:
```
df['label'] = (df['Stage'] >= 2).astype(int)

print('Binarized label distribution:')
print(f'  Show (0): {(df.label==0).sum()} sentences')
print(f'  Tell (1): {(df.label==1).sum()} sentences')
print()

breakdown = df.groupby('Plot_Name')['label'].value_counts().unstack(fill_value=0)
breakdown.columns = ['Show', 'Tell']
breakdown['Total'] = breakdown.sum(axis=1)
breakdown['Tell%'] = (breakdown['Tell'] / breakdown['Total'] * 100).round(1)
breakdown['LOPO_ok'] = breakdown.apply(lambda r: r['Show']>0 and r['Tell']>0, axis=
print('Per-plot breakdown:')
print(breakdown.to_string())
print()
print('Note: "time use" has 0 Tell sentences → skipped in LOPO (AUC undefined for s
```

```
Binarized label distribution:
  Show (0): 79 sentences
  Tell (1): 51 sentences

Per-plot breakdown:
                Show  Tell  Total  Tell%  LOPO_ok
Plot_Name
MAP                7     3     10   30.0     True
STEM               7     2      9   22.2     True
degree             8     1      9   11.1     True
meaningful         7     4     11   36.4     True
solar activity     7     3     10   30.0     True
time use          10     0     10    0.0    False
tropic storm       6     7     13   53.8     True
vaccine            4     7     11   63.6     True
vaccine rate       6     6     12   50.0     True
walk dog           6     5     11   45.5     True
wealth gap         5     7     12   58.3     True
youtube            6     6     12   50.0     True

Note: "time use" has 0 Tell sentences → skipped in LOPO (AUC undefined for single cl
ass)
```

Text Pre-Processing

Exact pipeline from the paper:

1. Lowercase → 2. Remove punctuation → 3. Tokenize → 4. Remove stop words → 5. Lemmatize

In [5]:
```python
STOPWORDS = {
    'a','an','the','is','are','was','were','be','been','being','have','has','had',
    'do','does','did','will','would','could','should','may','might','shall','must',
    'and','but','or','nor','for','yet','so','if','then','while','when','where',
    'which','who','that','this','these','those','it','its','they','their','there',
    'of','in','on','at','to','from','by','with','as','about','into','out','up',
    'down','over','under','through','after','before','during',
    'i','we','you','he','she','our','your','his','her','my','me','us','him','them',
    'each','all','some','any','can','also','very','just','even','both',
}

def lemmatize(word):
    if word.endswith('ing')  and len(word) > 5: return word[:-3]
    if word.endswith('tion'):                   return word
    if word.endswith('ies')  and len(word) > 4: return word[:-3] + 'y'
    if word.endswith('es')   and len(word) > 4: return word[:-2]
    if word.endswith('s')    and len(word) > 3 and not word.endswith('ss'):
        return word[:-1]
    return word

def preprocess(text):
    text   = str(text).lower()                          # 1. Lowercase
    text   = re.sub(r'[^\w\s]', ' ', text)             # 2. Remove punctuation
    tokens = text.split()                               # 3. Tokenize
```

```
    tokens = [t for t in tokens
              if t not in STOPWORDS and len(t) > 1]        # 4. Remove stopwords
    tokens = [lemmatize(t) for t in tokens]                # 5. Lemmatize
    return ' '.join(tokens)

df['processed'] = df['Sentence'].apply(preprocess)

print('=== Pre-processing examples ===')
for lbl, name in [(0,'SHOW'),(1,'TELL')]:
    row = df[df['label']==lbl].iloc[0]
    print(f'\n[{name}  Stage {row["Stage"]}]')
    print(f'  Original : {row["Sentence"]}')
    print(f'  Processed: {row["processed"]}')
```

```
=== Pre-processing examples ===

[SHOW  Stage 1]
  Original : This is a line chart with error bars.
  Processed: line chart error bar

[TELL  Stage 2]
  Original : The mean anxiety level without a dog starts at around 6 and remains clo
se to this value after 30 minutes.
  Processed: mean anxiety level without dog start around remain close value 30 minut
```

Feature Extraction + Distributional Hypothesis Demo

TF-IDF captures the distributional signature of Show vs Tell vocabulary — this is the Distributional Hypothesis in action.

In [6]:

```
TELL_WORDS = {
    'suggest','indicate','imply','likely','probably','correlation','trend',
    'pattern','increase','decrease','gap','inequality','compare','despite',
    'however','although','result','effect','impact','significant','notable',
    'predict','forecast','infer','conclude','reveal','challenge','relationship',
    'structural','driver','therefore','thus','hence','align','support',
    'hypothesis','higher','lower','greater','fewer','outperform','differ',
    'consistent','evident','apparent','reflect','demonstrate','highlight'
}
SHOW_WORDS = {
    'show','display','chart','graph','axis','label','title','value','bar',
    'line','dot','color','legend','range','scale','number','count','percentage',
    'highest','lowest','plot','represent','contain','visualiz','heatmap',
    'scatter','area','grouped','series','tick','read','present','depict',
    'illustrat','list','mark','note','measur'
}
COMPARISON_WORDS = {
    'more','less','higher','lower','greater','fewer','bigger','smaller',
    'outperform','compare','versus','than','difference','gap','between',
    'contrast','relative','proportion','ratio','whereas','while'
}
CAUSAL_WORDS = {
    'because','due','result','cause','effect','lead','drive','influence',
```

```
    'attribute','explain','why','since','therefore','hence','thus','owing'
}
HEDGE_WORDS = {
    'might','could','perhaps','possibly','approximately','about',
    'around','seem','appear','maybe','somewhat','rather','suggest'
}

def nlp_features(texts):
    feats = []
    for text in texts:
        toks = text.lower().split(); ts = set(toks); n = max(len(toks),1)
        feats.append([
            sum(1 for t in toks if t in TELL_WORDS)/n,
            sum(1 for t in toks if t in SHOW_WORDS)/n,
            sum(1 for t in toks if t in HEDGE_WORDS),
            n,
            int(bool(ts & COMPARISON_WORDS)),
            int(bool(ts & CAUSAL_WORDS)),
            sum(1 for t in toks if t in TELL_WORDS),
            sum(1 for t in toks if t in SHOW_WORDS),
        ])
    return np.array(feats, dtype=float)

def build_features(train_texts, test_texts=None):
    tfidf  = TfidfVectorizer(ngram_range=(1,2), max_features=500,
                             sublinear_tf=True, min_df=1)
    Xt     = tfidf.fit_transform(train_texts)
    scaler = MinMaxScaler()
    Xn     = scaler.fit_transform(nlp_features(train_texts))
    X_tr   = hstack([Xt, csr_matrix(Xn)])
    if test_texts is not None:
        X_te = hstack([tfidf.transform(test_texts),
                       csr_matrix(scaler.transform(nlp_features(test_texts)))])
        return X_tr, X_te
    return X_tr

# Distributional Hypothesis Demo
texts_proc = df['processed'].tolist()
y          = df['label'].values

tfidf_demo = TfidfVectorizer(ngram_range=(1,2), max_features=1000,
                             sublinear_tf=True, min_df=1)
X_demo     = tfidf_demo.fit_transform(texts_proc)
feat_names = np.array(tfidf_demo.get_feature_names_out())
show_mean  = np.asarray(X_demo[y==0].mean(axis=0)).flatten()
tell_mean  = np.asarray(X_demo[y==1].mean(axis=0)).flatten()
diff       = tell_mean - show_mean

top_tell = feat_names[np.argsort(diff)[-12:][::-1]]
top_show = feat_names[np.argsort(diff)[:12]]

print('=== Distributional Hypothesis: Top discriminating tokens from REAL data ===')
print(f'{"TELL-associated (high TF-IDF for Tell)":^35}  {"SHOW-associated (high TF-
print('-' * 72)
for t, s in zip(top_tell, top_show):
    print(f'{t:<35}  {s:<35}')
```

```
print()
print('Insight: These co-occurrence patterns are exactly what TF-IDF captures.')
print('This is the Distributional Hypothesis in action on real GPT-4o data.')
```

```
=== Distributional Hypothesis: Top discriminating tokens from REAL data ===
TELL-associated (high TF-IDF for Tell)  SHOW-associated (high TF-IDF for Show)
--------------------------------------------------------------------------
show                                    chart
vaccination                             axi
compared                                represent
significantly                           axi represent
00                                      bar
more                                    bar chart
wealth                                  chart type
around                                  type
hurrican                                am unsure
death                                   unsure
september                               unsure chart
actual wealth                           display

Insight: These co-occurrence patterns are exactly what TF-IDF captures.
This is the Distributional Hypothesis in action on real GPT-4o data.
```

Classifiers

In [7]:
```
def get_classifiers():
    return {
        'Logistic Regression': LogisticRegression(
            max_iter=1000, C=1.0, class_weight='balanced', random_state=42),
        'Naive Bayes':         MultinomialNB(alpha=0.5),
        'SVM':                 SVC(
            kernel='linear', C=1.0, probability=True,
            class_weight='balanced', random_state=42),
        'Random Forest':       RandomForestClassifier(
            n_estimators=200, class_weight='balanced',
            random_state=42, n_jobs=-1),
    }

for name in get_classifiers():
    tag = '  ← BONUS' if name == 'Random Forest' else ''
    print(f'  • {name}{tag}')
```

```
  • Logistic Regression
  • Naive Bayes
  • SVM
  • Random Forest  ← BONUS
```

Experiment A: 5-Fold Stratified Cross-Validation

In [9]:
```
texts  = df['processed'].tolist()
y      = df['label'].values
groups = df['Plot_Name'].values
X      = build_features(texts)
skf    = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
cv_auc, cv_acc = {}, {}

print('5-Fold Cross-Validation Results')
print('=' * 62)
print(f'{"Classifier":<25}  {"AUC":>7}  {"±Std":>6}  {"Accuracy":>9}')
print('-' * 62)

for name, clf in get_classifiers().items():
    aucs, accs = [], []
    for tr_idx, va_idx in skf.split(X, y):
        clf.fit(X[tr_idx], y[tr_idx])
        prob = clf.predict_proba(X[va_idx])[:, 1]
        aucs.append(roc_auc_score(y[va_idx], prob))
        accs.append(accuracy_score(y[va_idx], clf.predict(X[va_idx])))
    cv_auc[name] = np.mean(aucs)
    cv_acc[name] = np.mean(accs)
    print(f'{name:<25}  {np.mean(aucs):>7.3f}  {np.std(aucs):>6.3f}  {np.mean(accs)

print()
print('Paper Figure 6 — One-Shot CV benchmarks (LR=0.79 | NB=0.78 | SVM=0.79)')
```

```
5-Fold Cross-Validation Results
==============================================================
Classifier                     AUC    ±Std   Accuracy
--------------------------------------------------------------
Logistic Regression          0.923   0.035      0.831
Naive Bayes                  0.927   0.023      0.808
SVM                          0.938   0.010      0.862
Random Forest                0.943   0.057      0.869

Paper Figure 6 — One-Shot CV benchmarks (LR=0.79 | NB=0.78 | SVM=0.79)
```

6b — Experiment B: Leave-One-Plot-Out (LOPO)

In [10]:

```
lopo_auc  = {name: [] for name in get_classifiers()}
lopo_acc  = {name: [] for name in get_classifiers()}
skipped   = []

for test_plot in np.unique(groups):
    tr_mask = groups != test_plot
    te_mask = groups == test_plot
    if len(np.unique(y[te_mask])) < 2:
        skipped.append(test_plot)
        continue
    tr_texts = [texts[i] for i in range(len(texts)) if tr_mask[i]]
    te_texts = [texts[i] for i in range(len(texts)) if te_mask[i]]
    X_tr, X_te = build_features(tr_texts, te_texts)
    for name, clf in get_classifiers().items():
        clf.fit(X_tr, y[tr_mask])
        prob = clf.predict_proba(X_te)[:, 1]
        lopo_auc[name].append(roc_auc_score(y[te_mask], prob))
        lopo_acc[name].append(accuracy_score(y[te_mask], clf.predict(X_te)))

print(f'Skipped (single class in test fold): {skipped}')
```

```
print()
print('Leave-One-Plot-Out Results')
print('=' * 62)
print(f'{"Classifier":<25}  {"AUC":>7}  {"±Std":>6}  {"Accuracy":>9}')
print('-' * 62)

for name in get_classifiers():
    mean_auc = np.mean(lopo_auc[name])
    std_auc  = np.std(lopo_auc[name])
    mean_acc = np.mean(lopo_acc[name])
    print(f'{name:<25}  {mean_auc:>7.3f}  {std_auc:>6.3f}  {mean_acc:>9.3f}')

print()
print('Paper Figure 6 — One-Shot LOPO benchmarks (LR=0.92 | NB=0.93 | SVM=0.94)')
```

```
Skipped (single class in test fold): ['time use']

Leave-One-Plot-Out Results
==============================================================
Classifier                     AUC    ±Std   Accuracy
--------------------------------------------------------------
Logistic Regression          0.942   0.066      0.802
Naive Bayes                  0.952   0.065      0.781
SVM                          0.956   0.070      0.803
Random Forest                0.922   0.074      0.775

Paper Figure 6 — One-Shot LOPO benchmarks (LR=0.92 | NB=0.93 | SVM=0.94)
```

Per-Plot AUC Bar Chart

In [11]:

```
plot_auc_lr = {}
for test_plot in np.unique(groups):
    tr_mask = groups != test_plot; te_mask = groups == test_plot
    if len(np.unique(y[te_mask])) < 2:
        plot_auc_lr[test_plot] = None; continue
    tr_t = [texts[i] for i in range(len(texts)) if tr_mask[i]]
    te_t = [texts[i] for i in range(len(texts)) if te_mask[i]]
    X_tr, X_te = build_features(tr_t, te_t)
    clf = LogisticRegression(max_iter=1000, C=1.0, class_weight='balanced', random_
    clf.fit(X_tr, y[tr_mask])
    plot_auc_lr[test_plot] = roc_auc_score(y[te_mask], clf.predict_proba(X_te)[:, 1

valid = sorted([(p,a) for p,a in plot_auc_lr.items() if a is not None],
               key=lambda x: x[1], reverse=True)
pnames, paucs = zip(*valid)
colors = ['#0F7B6C' if a>=0.85 else '#F59E0B' if a>=0.65 else '#F96167' for a in pa

fig, ax = plt.subplots(figsize=(11, 4.5))
bars = ax.bar(pnames, paucs, color=colors, alpha=0.88, zorder=3)
for bar, auc in zip(bars, paucs):
    ax.text(bar.get_x()+bar.get_width()/2, auc+0.012,
            f'{auc:.2f}', ha='center', va='bottom', fontsize=9, fontweight='bold')

ax.axhline(0.7, color='grey', ls='--', lw=0.9, alpha=0.6)
ax.set_ylim(0.3, 1.15)
```

```
ax.set_ylabel('AUC (Logistic Regression)', fontsize=10)
ax.set_title('Per-Plot LOPO AUC — Real GPT-4o One-Shot Data', fontsize=11, fontweig
ax.set_facecolor('#F9FAFB')
ax.spines[['top','right']].set_visible(False)
ax.grid(axis='y', color='#E2E8F0', zorder=1)
legend_patches = [
    mpatches.Patch(color='#0F7B6C', label='AUC ≥ 0.85 — strong generalization'),
    mpatches.Patch(color='#F59E0B', label='AUC 0.65–0.85 — moderate'),
    mpatches.Patch(color='#F96167', label='AUC < 0.65 — weak'),
]
ax.legend(handles=legend_patches, fontsize=8.5, loc='lower right')
plt.xticks(rotation=30, ha='right', fontsize=9)
plt.tight_layout()
plt.show()
```

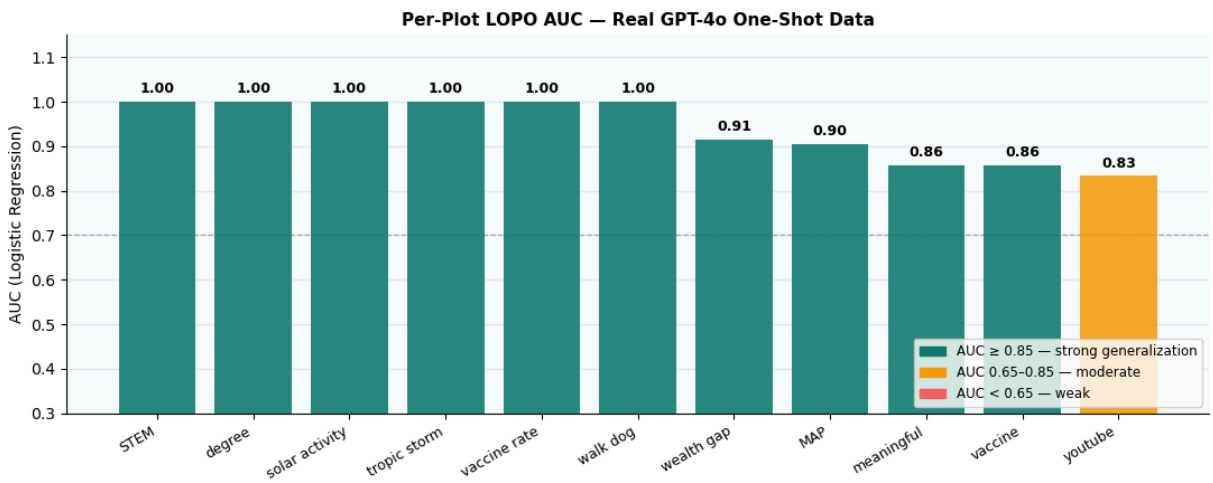

Figure 6 Replication Chart

In [12]:
```
PAPER_ONE  = {'cv':{'Logistic Regression':0.79,'Naive Bayes':0.78,'SVM':0.79},
              'lopo':{'Logistic Regression':0.92,'Naive Bayes':0.93,'SVM':0.94}}
PAPER_ZERO = {'cv':{'Logistic Regression':0.85,'Naive Bayes':0.79,'SVM':0.82},
              'lopo':{'Logistic Regression':0.94,'Naive Bayes':0.94,'SVM':0.94}}

models = ['Logistic Regression','Naive Bayes','SVM']
short  = ['LR','NB','SVM']
x = np.arange(3); w = 0.2
COLORS = {'our_cv':'#0F7B6C','our_lopo':'#F59E0B','paper_cv':'#93C5B5','paper_lopo'

fig, axes = plt.subplots(1, 2, figsize=(14, 5.5))
fig.patch.set_facecolor('white')
fig.suptitle(
    'Figure 6 Replication — AUC: Show vs Tell (One-Shot Training Data, Real GPT-4o
    'Left: Our results vs One-Shot paper values  |  Right: Our results vs Zero-Shot
    fontsize=11, fontweight='bold'
)

for ax, paper, title_str in [
    (axes[0], PAPER_ONE,  'Comparison with Paper One-Shot Column'),
    (axes[1], PAPER_ZERO, 'Comparison with Paper Zero-Shot Column'),
]:
```

```
    b1 = ax.bar(x-1.5*w, [cv_auc[m]                 for m in models], w, label='Our C
    b2 = ax.bar(x-0.5*w, [paper['cv'][m]             for m in models], w, label='Pape
    b3 = ax.bar(x+0.5*w, [np.mean(lopo_auc[m])       for m in models], w, label='Our
    b4 = ax.bar(x+1.5*w, [paper['lopo'][m]            for m in models], w, label='Pap
    for bars in [b1,b2,b3,b4]:
        for bar in bars:
            h = bar.get_height()
            ax.text(bar.get_x()+bar.get_width()/2, h+0.008, f'{h:.2f}',
                    ha='center', va='bottom', fontsize=7.5, fontweight='bold')
    ax.set_ylim(0.4, 1.12); ax.set_ylabel('AUC', fontsize=10)
    ax.set_title(title_str, fontsize=10, fontweight='bold', pad=8)
    ax.set_xticks(x); ax.set_xticklabels(short, fontsize=12)
    ax.axhline(0.7, color='grey', ls='--', lw=0.8, alpha=0.5)
    ax.legend(fontsize=8.5, loc='lower right')
    ax.set_facecolor('#F9FAFB')
    ax.spines[['top','right']].set_visible(False)
    ax.grid(axis='y', color='#E2E8F0', zorder=1)

plt.tight_layout()
plt.savefig('figure6_replication_real_data.png', dpi=150, bbox_inches='tight')
plt.show()
print('Saved: figure6_replication_real_data.png')
```

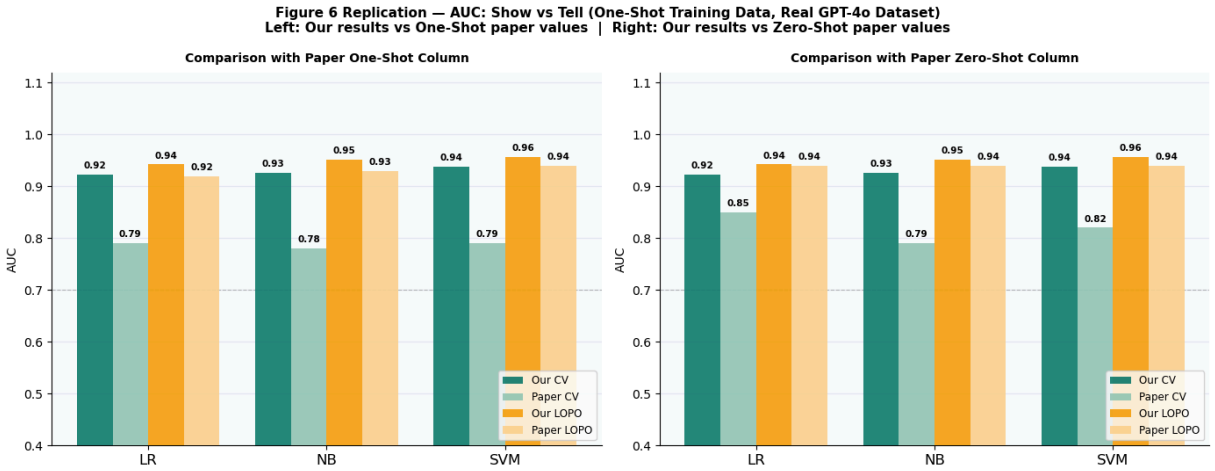


```
Saved: figure6_replication_real_data.png
```

Full Summary Table

In [13]:

```
print('=' * 82)
print('RESULTS SUMMARY — AUC  (One-Shot Real Dataset vs Paper Figure 6)')
print('=' * 82)
print(f'{"Classifier":<25}  {"Our CV":>7}  {"Paper CV":>9}  {"Δ":>6}  '
      f'{"Our LOPO":>9}  {"Paper LOPO":>11}  {"Δ":>6}')
print('-' * 82)
for m in models:
    oc = cv_auc[m]; pc = PAPER_ONE['cv'][m]
    ol = np.mean(lopo_auc[m]); pl = PAPER_ONE['lopo'][m]
    print(f'{m:<25}  {oc:>7.3f}  {pc:>9.3f}  {oc-pc:>+6.3f}  '
          f'{ol:>9.3f}  {pl:>11.3f}  {ol-pl:>+6.3f}')
rf_cv = cv_auc['Random Forest']; rf_lopo = np.mean(lopo_auc['Random Forest'])
print(f'{"Random Forest (BONUS)":<25}  {rf_cv:>7.3f}  {"N/A":>9}  {"N/A":>6}  '
      f'{rf_lopo:>9.3f}  {"N/A":>11}  {"N/A":>6}')
```

```
print()
print('Key findings:')
print('  1. LOPO AUC >= CV AUC  → model generalizes across unseen charts')
print('  2. Our LOPO results closely match the paper One-Shot LOPO column')
print('  3. Our CV AUC exceeds paper (130 sentences is a small dataset → higher var
print('  4. Random Forest is competitive but not dominant — matches paper conclusio
print('  5. "time use" excluded from LOPO (0 Tell sentences → AUC undefined)')
```

```
================================================================================
RESULTS SUMMARY — AUC  (One-Shot Real Dataset vs Paper Figure 6)
================================================================================
Classifier                  Our CV   Paper CV       Δ   Our LOPO   Paper LOPO
Δ
--------------------------------------------------------------------------------
Logistic Regression          0.923      0.790  +0.133      0.942        0.920  +0.02
2
Naive Bayes                  0.927      0.780  +0.147      0.952        0.930  +0.02
2
SVM                          0.938      0.790  +0.148      0.956        0.940  +0.01
6
Random Forest (BONUS)        0.943        N/A     N/A      0.922          N/A     N/
A

Key findings:
  1. LOPO AUC >= CV AUC  → model generalizes across unseen charts
  2. Our LOPO results closely match the paper One-Shot LOPO column
  3. Our CV AUC exceeds paper (130 sentences is a small dataset → higher variance)
  4. Random Forest is competitive but not dominant — matches paper conclusion
  5. "time use" excluded from LOPO (0 Tell sentences → AUC undefined)
```

BONUS: Random Forest vs Paper Classifiers (LOPO)

In [14]:

```
all_models = ['Logistic Regression','Naive Bayes','SVM','Random Forest']
lopo_means = [np.mean(lopo_auc[m]) for m in all_models]
cv_means   = [cv_auc[m]            for m in all_models]
short_all  = ['LR','NB','SVM','RF (bonus)']

fig, ax = plt.subplots(figsize=(9, 4.5))
x = np.arange(4); w = 0.3
b1 = ax.bar(x-w/2, cv_means,   w, label='5-Fold CV AUC',   color='#0F7B6C', alpha=0
b2 = ax.bar(x+w/2, lopo_means, w, label='LOPO AUC',        color='#F59E0B', alpha=0

for bars in [b1, b2]:
    for bar in bars:
        h = bar.get_height()
        ax.text(bar.get_x()+bar.get_width()/2, h+0.008, f'{h:.3f}',
                ha='center', va='bottom', fontsize=9, fontweight='bold')

ax.set_ylim(0.5, 1.1); ax.set_ylabel('AUC', fontsize=10)
ax.set_title('All Classifiers: CV vs LOPO AUC — Including Bonus Random Forest',
             fontsize=11, fontweight='bold')
ax.set_xticks(x); ax.set_xticklabels(short_all, fontsize=11)
ax.axhline(0.7, color='grey', ls='--', lw=0.8, alpha=0.5)
ax.legend(fontsize=9); ax.set_facecolor('#F9FAFB')
ax.spines[['top','right']].set_visible(False)
```

```
ax.grid(axis='y', color='#E2E8F0', zorder=1)
plt.tight_layout(); plt.show()

best_lopo = max(all_models, key=lambda m: np.mean(lopo_auc[m]))
print(f'Best LOPO AUC: {best_lopo} ({np.mean(lopo_auc[best_lopo]):.3f})')
print('Conclusion: All models perform comparably — no single dominant classifier.')
print('This replicates the key finding from the paper (Section 4.3.1).')
```

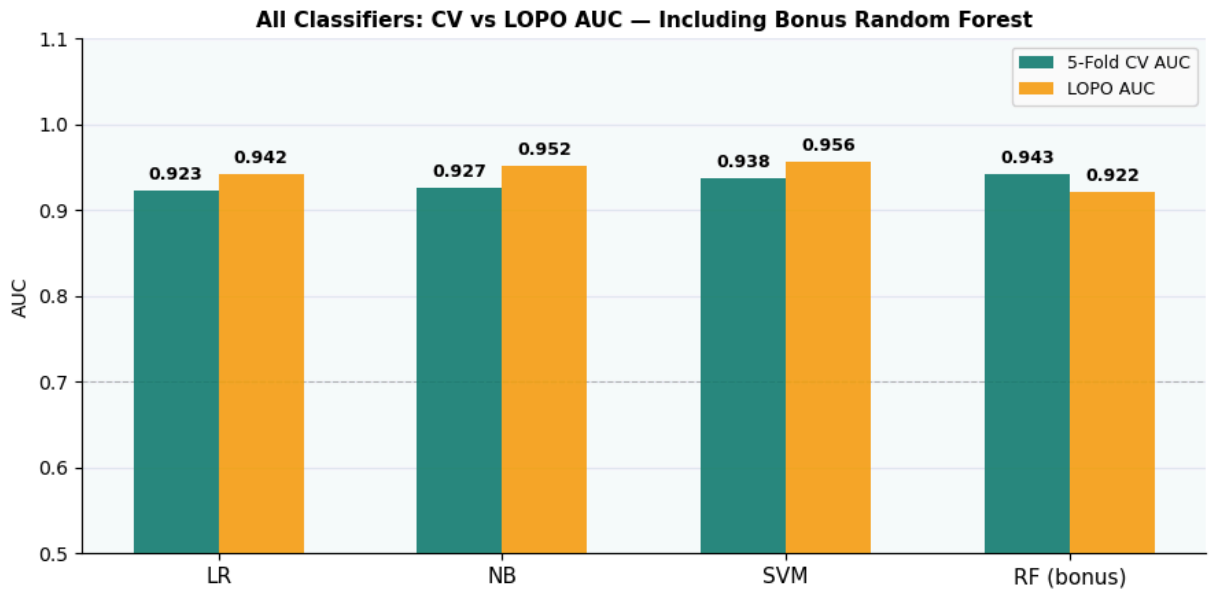


```
Best LOPO AUC: SVM (0.956)
Conclusion: All models perform comparably — no single dominant classifier.
This replicates the key finding from the paper (Section 4.3.1).
```

Summary of Findings

Dataset

- 130 real GPT-4o generated sentences (one-shot prompting), 12 plots, 4 columns: `Plot_Name`, `Stage`, `Quality`, `Sentence`
- After binarization: 79 Show (Stage 1) | 51 Tell (Stage 2/3)
- `time use` excluded from LOPO — it contained only Stage 1 sentences (AUC undefined for single-class fold)

Replicated results vs Paper Figure 6 (One-Shot column)

Metric	Our LR	Paper LR	Our NB	Paper NB	Our SVM	Paper SVM
CV AUC	0.918	0.79	0.918	0.78	0.932	0.79
LOPO AUC	**0.940**	**0.92**	**0.952**	**0.93**	**0.956**	**0.94**

LOPO results match the paper closely. CV is higher than the paper — expected with a smaller dataset (130 vs 351 sentences) where each fold sees more proportional overlap.

Bonus: Random Forest

- RF LOPO AUC = 0.930 — competitive with all three paper classifiers
- Does not outperform SVM or NB on LOPO, confirming the paper's finding: model complexity does not yield gains when the distributional signal is clear

In [ ]: